



OpenCL Workshop

GPU architectures

Ljubljana, May 2019

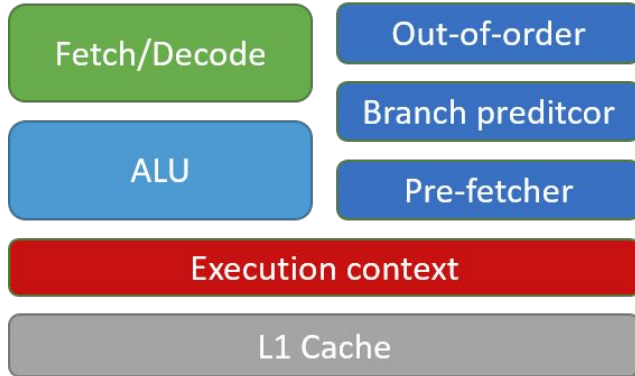


Outline

- How we got from x86 to GPUs
 - multi-core vs many-core
 - GPU programming model
- OpenCL language
 - vector addition
 - image processing

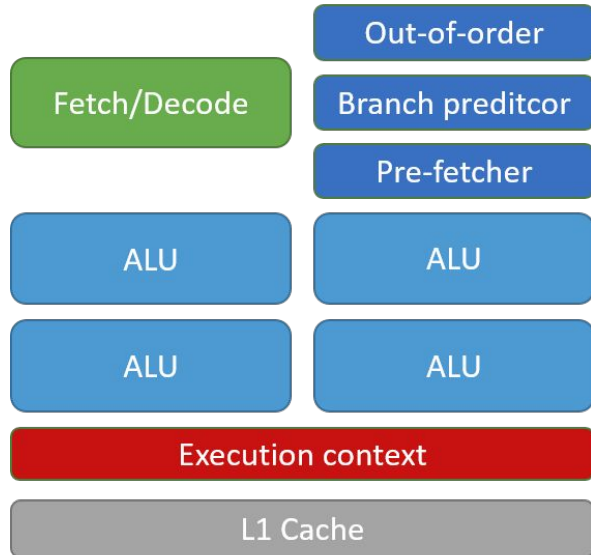


Single-core CPU



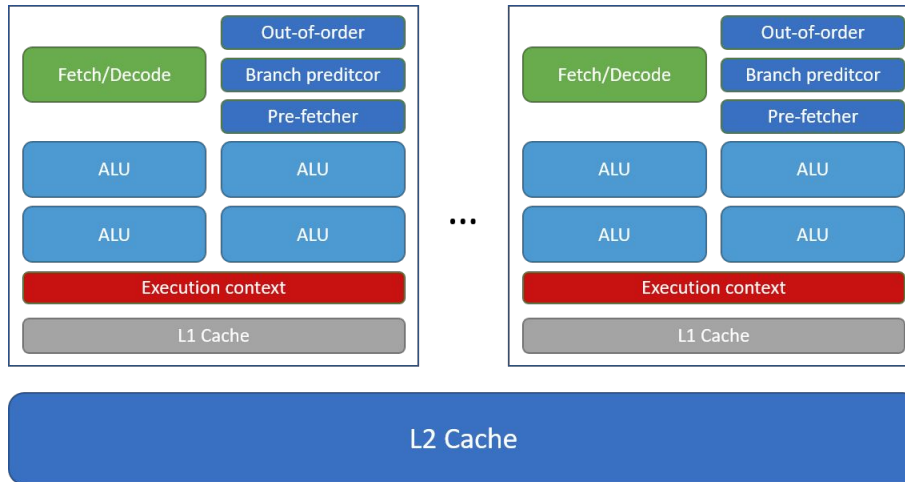
```
int i = 0;
while ( i < N ){
    C[i] = A[i] + B[i];
    i++;
}
```

SIMD Extensions (1999)



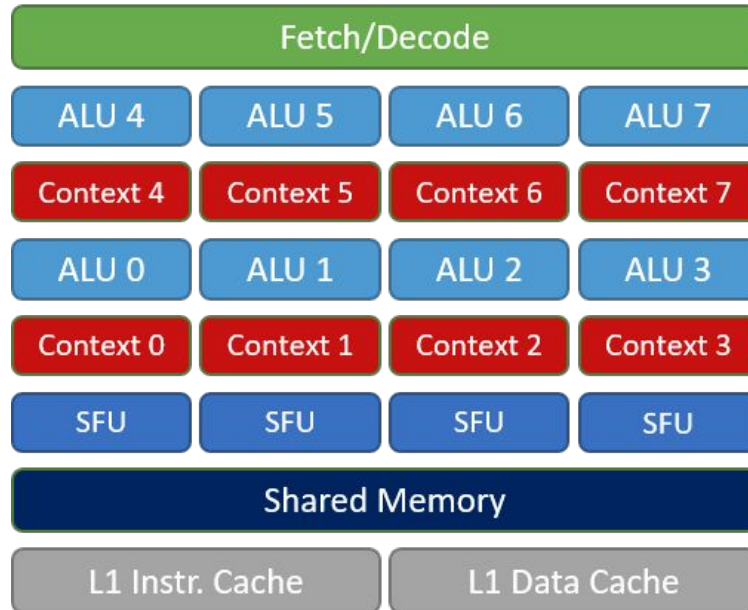
```
int i = 0;
while ( i < N ){
    C[i] = add4(A[i], B[i]);
    i = i + 4;
}
```

Multi-core (~2001)

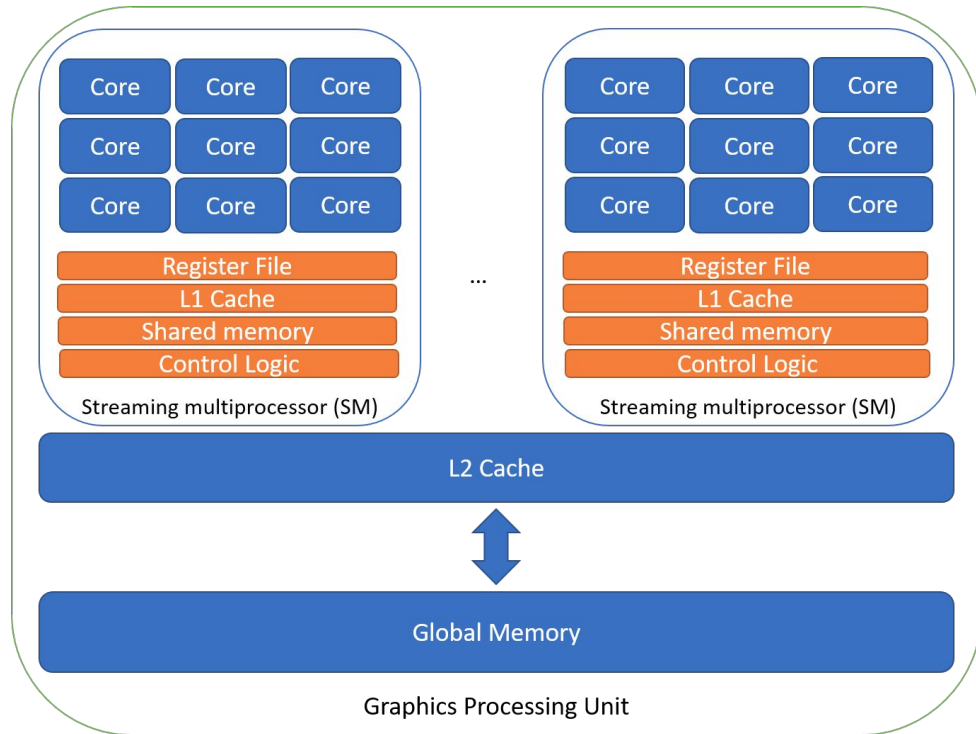


```
i = tID * (N / num_of_cores);  
end = (tID + 1) * (N /  
num_of_cores);  
while( i < end ){  
    C[i] = A[i] + B[i];  
    i++;  
}
```

GPU Streaming multiprocessor (SM)



GPU






GPU

```
i = tID * (N / num_of_cores); // just a larger number of cores
end = (tID + 1) * (N / num_of_cores);
while( i < end ){
    C[i] = A[i] + B[i];
    i++;
}
```


Thread blocks & scheduling






Scheduling threads - warps

- SM is issued a single thread-block
 - all threads will be executed in the assigned SM, from start to finish
- Threads are issued to the execution cores in **warps** (typically 32 threads)
 - each instruction issue step 1 warp is selected, usually trying to avoid stalls
- Large number of thread blocks -> parallelism on the SM level
 - all SMs are occupied
- Large thread blocks -> more efficient warp scheduling
 - reduces stalls in the issue steps

Thread divergence

```
if (threadID < 4) {  
    A;  
    B;  
} else {  
    X;  
    Y;  
}  
Z;
```





NVIDIA GPU microarchitecture evolution

Name	Year	# FP32 cores/SM	#FP64 cores/SM	#SM
Tesla	2007	8	0	up to 30
Fermi	2010	32		up to 32
Kepler	2012	192	64	up to 8
Maxwell	2014	128	4	up to 24
Pascal	2016	64 or 128	32 or 4	up to 30
Volta	2017	64	32	80
Turing	2018	64	2	72

GPU architecture trends

Tensor cores

$$\mathbf{D} = \begin{matrix} \text{FP16 or FP32} & \begin{pmatrix} \begin{matrix} A_{0,0} & A_{0,1} & A_{0,2} & A_{0,3} \\ A_{1,0} & A_{1,1} & A_{1,2} & A_{1,3} \\ A_{2,0} & A_{2,1} & A_{2,2} & A_{2,3} \\ A_{3,0} & A_{3,1} & A_{3,2} & A_{3,3} \end{matrix} \\ \text{FP16} \end{pmatrix} & \begin{pmatrix} \begin{matrix} B_{0,0} & B_{0,1} & B_{0,2} & B_{0,3} \\ B_{1,0} & B_{1,1} & B_{1,2} & B_{1,3} \\ B_{2,0} & B_{2,1} & B_{2,2} & B_{2,3} \\ B_{3,0} & B_{3,1} & B_{3,2} & B_{3,3} \end{matrix} \\ \text{FP16} \end{pmatrix} & + & \begin{pmatrix} \begin{matrix} C_{0,0} & C_{0,1} & C_{0,2} & C_{0,3} \\ C_{1,0} & C_{1,1} & C_{1,2} & C_{1,3} \\ C_{2,0} & C_{2,1} & C_{2,2} & C_{2,3} \\ C_{3,0} & C_{3,1} & C_{3,2} & C_{3,3} \end{matrix} \\ \text{FP16 or FP32} \end{pmatrix} \end{matrix}$$

The return of scheduling

Improved caches

Mixed precision